

MLP_PYTORCH2

September 9, 2023

Necessary Imports

```
[2]: import torch
import numpy as np
from torchvision import datasets
import torchvision.transforms as transforms
from torch.utils.data.sampler import SubsetRandomSampler
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
from sklearn.metrics import confusion_matrix
import matplotlib.pyplot as plt
import seaborn as sns

!pip install -Uqq ipdb
import ipdb
```

1.6/1.6 MB

7.8 MB/s eta 0:00:00

Selecting the device

```
[3]: if torch.cuda.is_available():
    device = torch.device('cuda')
else:
    device = torch.device('cpu')
device = torch.device('cpu')
```

The Neural Network

```
[4]: # define NN architecture
class Net(nn.Module):
    def __init__(self, input_dim, output_dim, hidden_layer_dimensions, dropout):

        super(Net, self).__init__()

        # number of hidden nodes in each layer (512)
```

```

        self.dropout = nn.Dropout(dropout)
        self.input_dim = input_dim
        self.output_dim = output_dim

        self.d_layer = [input_dim] + hidden_layer_dimensions + [output_dim]
        layer_list = [nn.Linear(self.d_layer[l], self.d_layer[l+1]) for l in
↪range(len(self.d_layer) - 1)]
        self.linears = nn.ModuleList(layer_list)

    def forward(self,X):

        X = X.view(-1, self.input_dim)
        # relu(Wl x) for all hidden layer
        for layer in self.linears[:-1]:
            X = F.relu(layer(X))
            X = self.dropout(X)
        # softmax(Wl x) for output layer
        out = self.linears[-1](X)
        return F.log_softmax(out, dim=1)

```

Class to perform operations on the neural network

```

[5]: %pdb on
class Classifier:

    #def __init__(self):
    #    model = Net()

    def load_data(self, batch_size = 64):
        # prepare data loaders

        # how many samples per batch to load
        self.batch_size = batch_size
        # percentage of training set to use as validation
        valid_size = 0.2
        # convert data to torch.FloatTensor
        transform = transforms.ToTensor()

        # choose the training and testing datasets
        train_validation_data = datasets.MNIST(root = 'data', train = True,
↪download = True, transform = transform)
        self.test_data = datasets.MNIST(root = 'data', train = False, download =
↪True, transform = transform)

```

```

# obtain training indices that will be used for validation
num_train_validation = len(train_validation_data)
self.num_test = len(self.test_data)

indices = list(range(num_train_validation))
np.random.shuffle(indices)
split = int(np.floor(valid_size * num_train_validation))

train_index, valid_index = indices[split:], indices[:split]

# define samplers for obtaining training and validation batches
train_sampler = SubsetRandomSampler(train_index)
valid_sampler = SubsetRandomSampler(valid_index)

# prepare data loaders
self.train_loader = torch.utils.data.DataLoader(train_validation_data,
↪batch_size = batch_size,
                                                    sampler = train_sampler,
↪num_workers = 0)
self.valid_loader = torch.utils.data.DataLoader(train_validation_data,
↪batch_size = batch_size,
                                                    sampler = valid_sampler,
↪num_workers = 0)
self.test_loader = torch.utils.data.DataLoader(self.test_data, batch_size =
↪batch_size,
                                                    num_workers = 0)

self.image_height = self.test_data[0][0].shape[1]
self.image_width = self.test_data[0][0].shape[2]

self.input_dim = self.image_height * self.image_width

self.output_dim = np.unique(self.test_data.targets).shape[0]

print("data loaded")

def create_model(self, opt = "adam", weight_decay = 0.001, learning_rate = 0.
↪01, beta1 = 0.9, beta2 = 0.999, epsilon = 1e-8, hidden_layer_dimensions =
↪[500, 250, 100], dropout = 0.2):

```

```

        self.model = Net(self.input_dim, self.output_dim, hidden_layer_dimensions,
↳dropout).to(device)
        self.optimizers = {"sgd" : optim.SGD(self.model.parameters(), lr =
↳learning_rate, momentum = 0.5, weight_decay = weight_decay),
                            "adam" : optim.Adam(self.model.parameters(), lr =
↳learning_rate, betas = [beta1, beta2], eps = epsilon)}
        self.optimizer = self.optimizers[opt]

        print(self.model)

def train_model(self, criterion = nn.NLLLoss() ,epochs = 15, log_interval =
↳200):

    losses, accuracies = dict(train = [], val = []), dict(train=[], val = [])

    for epoch in range(epochs):

        #train phase
        self.model.train() #set the model in training mode
        train_loss = 0.0
        train_corrects = 0
        total = 0

        for batchTrain, (images, labels) in enumerate(self.train_loader):
            train_loss = 0.0
            train_corrects = 0

            images = images.to(device)
            labels = labels.to(device)
            total = len(labels)
            # print(images.shape)

            self.optimizer.zero_grad()

            with torch.set_grad_enabled(True):
                outputs = self.model(images)
                _, preds = torch.max(outputs, 1)
                # print(outputs.device)
                # print(labels.device)
                # print(outputs)
                loss = criterion(outputs, labels)

                # backward + optimize only if in training phase
                loss.backward()
                self.optimizer.step()

```

```

        train_loss += loss.item() * images.size(0)
        train_corrects += torch.sum(preds == labels.data)

    if(batchTrain % log_interval == 0):
        losses["train"].append(train_loss/total)
        accuracies["train"].append(train_corrects/total)

    validation_loss = 0
    validation_corrects = 0
    total = 0

    for batchValidation, (images, labels) in enumerate(self.valid_loader):

        images = images.to(device)
        labels = labels.to(device)

        total += len(images)
        self.model.eval()

        with torch.set_grad_enabled(False):
            outputs = self.model(images)
            _, preds = torch.max(outputs, 1)
            loss = criterion(outputs, labels)

            validation_loss += loss.item() * images.size(0)
            validation_corrects += torch.sum(preds == labels.data)

        if(batchTrain % log_interval == 0):
            losses["val"].append(validation_loss/total)
            accuracies["val"].append(validation_corrects/total)

    return losses, accuracies

def plotConfusionMatrix(self, yTrue, yPred):
    # Create a heatmap of the confusion matrix
    confusion_mat = confusion_matrix(yTrue, yPred)
    plt.figure(figsize=(8, 6))
    sns.heatmap(confusion_mat, annot=True, fmt="d", cmap="Blues")
    plt.xlabel("Predicted")
    plt.ylabel("True")
    plt.title("Confusion Matrix")
    plt.show()

def test_model(self, criterion = nn.NLLLoss(), log_interval = 200):

```

```

self.model.eval()
test_loss = 0
test_corrects = 0
yPred = []
yTrue = []
total = 0

for batchTest, (images, labels) in enumerate(self.test_loader):

    images = images.to(device)
    labels = labels.to(device)
    yTrue += labels

    total += len(images)

    with torch.set_grad_enabled(False):
        outputs = self.model(images)
        _, preds = torch.max(outputs, 1)
        yPred += preds
        loss = criterion(outputs, labels)

    test_loss += loss.item() * images.size(0)
    test_corrects += torch.sum(preds == labels.data)

self.plotConfusionMatrix(yTrue, yPred)
return (test_loss/total, test_corrects/total)

```

Automatic pdb calling has been turned ON

```

[6]: def experiment(exp_id, batch_size = 64, opt = "sgd", weight_decay = 0.0,
    ↪ hidden_layer_dimensions = [500, 250, 100], dropout = 0.0, epochs = 15,
    ↪ log_interval = 200):

    print("experiment ", exp_id, "started")
    c = Classifier()
    c.load_data(batch_size = batch_size)
    c.create_model(opt = opt, weight_decay = weight_decay,
    ↪ hidden_layer_dimensions = hidden_layer_dimensions, dropout = dropout)
    (trainLosses, accuracies) = c.train_model(epochs = epochs, log_interval =
    ↪ log_interval)

    (test_loss, test_accuracy) = c.test_model()

    plt.title("Train & validation losses vs epochs")
    plt.plot(np.linspace(0, len(trainLosses["train"])-1,
    ↪ len(trainLosses["train"])), trainLosses["train"], label = "train losses")

```

```

plt.plot(np.linspace(0, len(trainLosses["val"])-1, len(trainLosses["val"])),
↪trainLosses["val"], label = "validation losses")
plt.show()
plt.title("Train & validation accuracies vs epochs")
plt.plot(np.linspace(0, len(accuracies["train"])-1,
↪len(accuracies["train"]))), accuracies["train"], label = "train accuracies")
plt.plot(np.linspace(0, len(accuracies["val"])-1, len(accuracies["val"])),
↪accuracies["val"], label = "validation accuracies")
plt.legend()
plt.show()

print("test loss is: ", test_loss)
print("test accuracy is: ", test_accuracy)
print("experiment: ", exp_id, " completed")

```

Experimenting the performance with different parameters.

Experiment: With best performing configuration in The Neural Network written from scratch,
Obtaining an accuracy of 93% in 5 epochs

```

[7]: #experiment : best performance from the Neural Network written from scratch
↪mimicked
experiment(exp_id = 1, batch_size = 64,opt = "sgd", weight_decay = 0.0,
↪hidden_layer_dimensions = [500, 250, 100], dropout = 0.0, epochs = 5,
↪log_interval = 200)

```

experiment 1 started

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/train-images-idx3-ubyte.gz> to
data/MNIST/raw/train-images-idx3-ubyte.gz

100%| | 9912422/9912422 [00:00<00:00, 72603701.89it/s]

Extracting data/MNIST/raw/train-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/train-labels-idx1-ubyte.gz> to
data/MNIST/raw/train-labels-idx1-ubyte.gz

100%| | 28881/28881 [00:00<00:00, 26588168.09it/s]

Extracting data/MNIST/raw/train-labels-idx1-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz>

Downloading <http://yann.lecun.com/exdb/mnist/t10k-images-idx3-ubyte.gz> to
data/MNIST/raw/t10k-images-idx3-ubyte.gz

100%| | 1648877/1648877 [00:00<00:00, 23077818.44it/s]

Extracting data/MNIST/raw/t10k-images-idx3-ubyte.gz to data/MNIST/raw

Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz>
 Downloading <http://yann.lecun.com/exdb/mnist/t10k-labels-idx1-ubyte.gz> to
 data/MNIST/raw/t10k-labels-idx1-ubyte.gz

100%| | 4542/4542 [00:00<00:00, 13322048.09it/s]

Extracting data/MNIST/raw/t10k-labels-idx1-ubyte.gz to data/MNIST/raw

data loaded

Net(

(dropout): Dropout(p=0.0, inplace=False)

(linears): ModuleList(

(0): Linear(in_features=784, out_features=500, bias=True)


(1): Linear(in_features=500, out_features=250, bias=True)


(2): Linear(in_features=250, out_features=100, bias=True)


(3): Linear(in_features=100, out_features=10, bias=True)

)

)







test loss is: 0.22002335015535354

test accuracy is: tensor(0.9338)

experiment: 1 completed

Using Adam optimizer: the obtained accuracy is 94.6% on test data


```
[ ]: #experiment : using adam optimizer
experiment(exp_id = 2, batch_size = 64, opt = "adam", weight_decay = 0.0,
↳hidden_layer_dimensions = [500, 250, 100], dropout = 0.0, epochs = 15,
↳log_interval = 200)
```


experiment 2 started


data loaded

```
Net(
  (dropout): Dropout(p=0.0, inplace=False)
  (linears): ModuleList(
    (0): Linear(in_features=784, out_features=500, bias=True)
    (1): Linear(in_features=500, out_features=250, bias=True)
    (2): Linear(in_features=250, out_features=100, bias=True)
    (3): Linear(in_features=100, out_features=10, bias=True)
  )
)
```

1
0







test loss is: 0.19614677001610398

test accuracy is: tensor(0.9466)

experiment: 2 completed

Experiment: Using SGD as optimizer, and L2 regularizer Obtained Accuracy: 64%

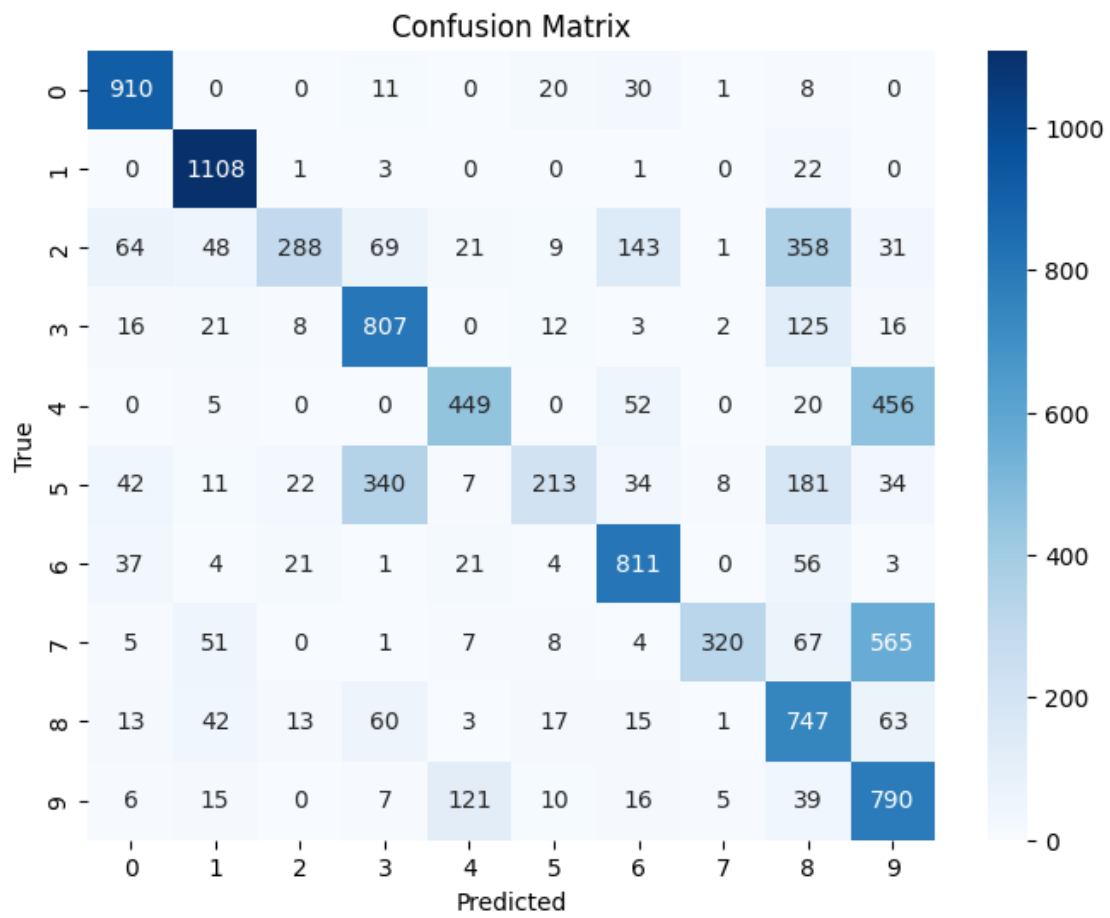
```
[ ]: #experiment : using l2 regularizer
experiment(exp_id = 3, batch_size = 64, opt = "sgd", weight_decay = 0.01,
↳hidden_layer_dimensions = [500, 250, 100], dropout = 0.0, epochs = 15,
↳log_interval = 200)
```


experiment 3 started


data loaded

```
Net(
  (dropout): Dropout(p=0.0, inplace=False)
  (linears): ModuleList(
    (0): Linear(in_features=784, out_features=500, bias=True)
    (1): Linear(in_features=500, out_features=250, bias=True)
    (2): Linear(in_features=250, out_features=100, bias=True)
    (3): Linear(in_features=100, out_features=10, bias=True)
  )
)
```

1
0







test loss is: 1.1016543467521667

test accuracy is: tensor(0.6443)

experiment: 3 completed

Experiment: Using SGD as optimizer and dropout regularizer, Obtained Accuracy: 76%


```
[ ]: #experiment : using dropout
experiment(exp_id = 4, batch_size = 64, opt = "sgd", weight_decay = 0.00,
↳hidden_layer_dimensions = [500, 250, 100], dropout = 0.02, epochs = 15,
↳log_interval = 200)
```


experiment 4 started


data loaded

```
Net(
  (dropout): Dropout(p=0.02, inplace=False)
  (linears): ModuleList(
    (0): Linear(in_features=784, out_features=500, bias=True)
    (1): Linear(in_features=500, out_features=250, bias=True)
    (2): Linear(in_features=250, out_features=100, bias=True)
    (3): Linear(in_features=100, out_features=10, bias=True)
  )
)
```


1
0







test loss is: 0.7571557574272155
test accuracy is: tensor(0.7672)
experiment: 4 completed

Experiment: Using sgd as optimizer and both dropout and l2 regularizer, Obtained accuracy: 74%

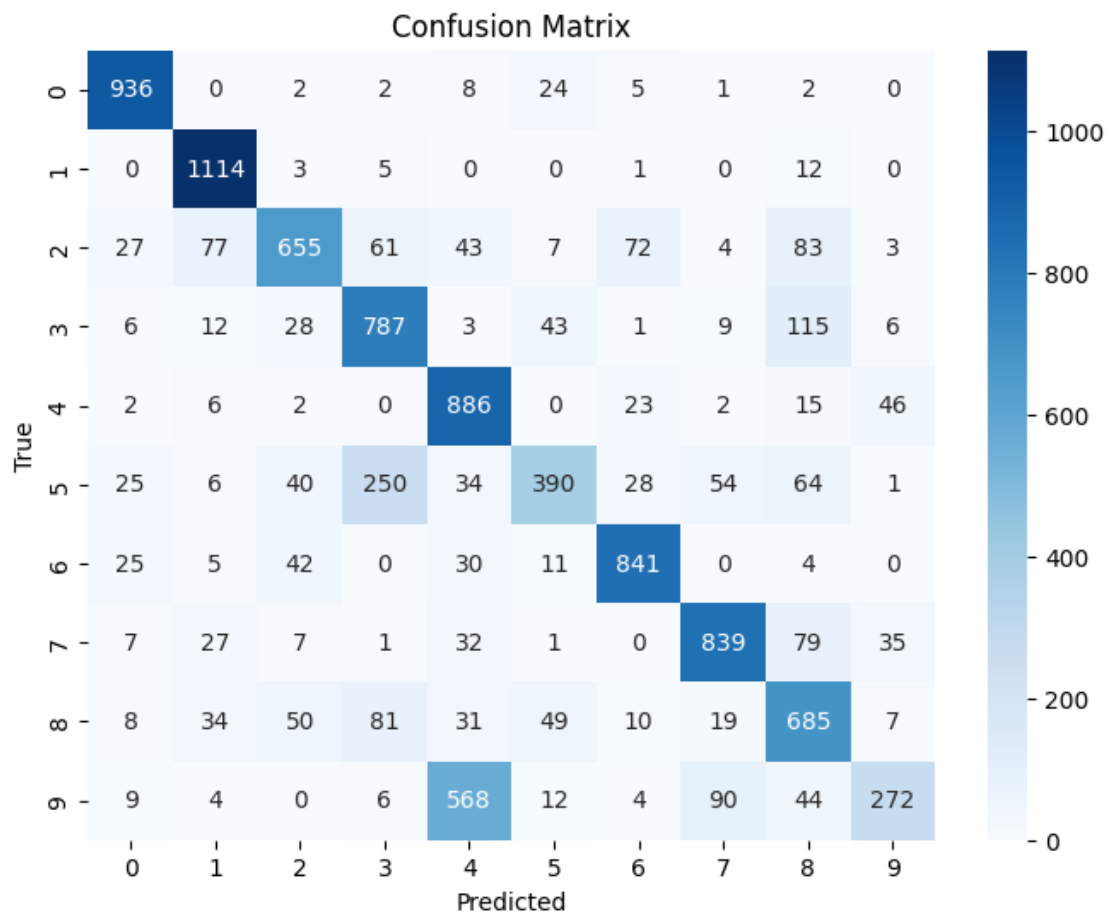
```
[ ]: #experiment : using both
experiment(exp_id = 5, batch_size = 64, opt = "sgd", weight_decay = 0.01,
↳hidden_layer_dimensions = [500, 250, 100], dropout = 0.02, epochs = 1,
↳log_interval = 200)
```

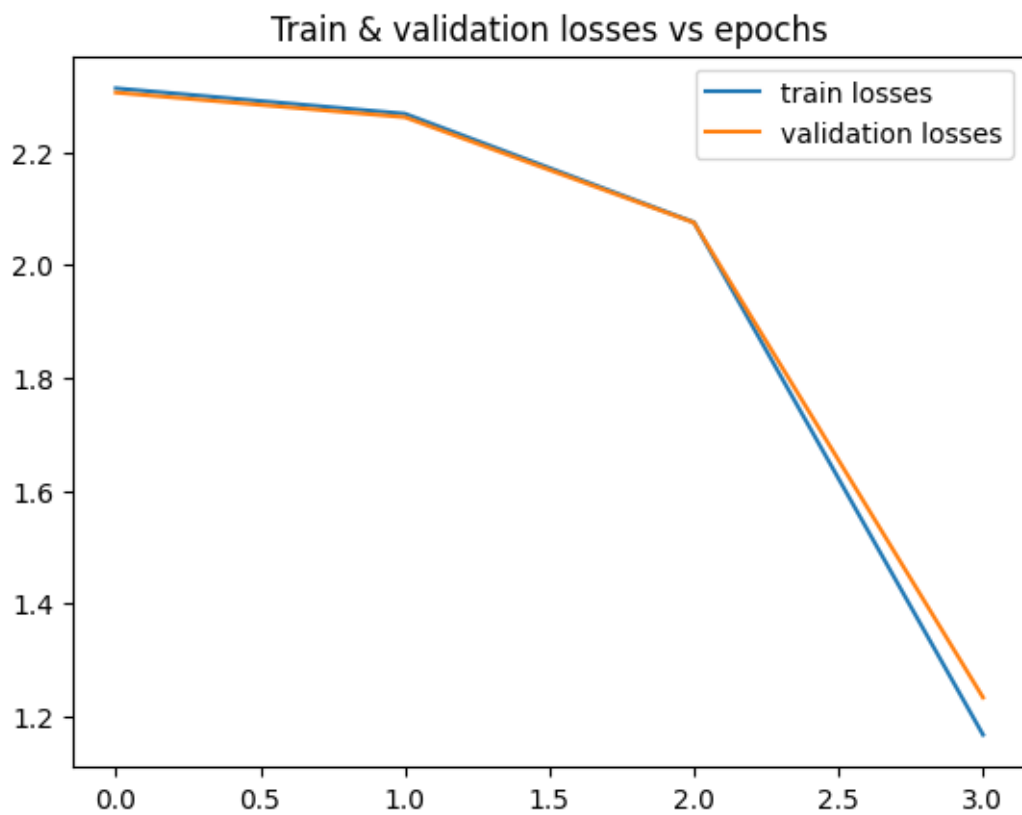
experiment 5 started


data loaded

```
Net(
  (dropout): Dropout(p=0.02, inplace=False)
  (linears): ModuleList(
    (0): Linear(in_features=784, out_features=500, bias=True)
    (1): Linear(in_features=500, out_features=250, bias=True)
    (2): Linear(in_features=250, out_features=100, bias=True)
    (3): Linear(in_features=100, out_features=10, bias=True)
  )
)
```

1
0







```
test loss is: 0.8179346004486084
test accuracy is: tensor(0.7405)
experiment: 5 completed
```

Observation: The SGD and Adam performs realitively very similar i.e the accuracy difference is lower. But performance with use of regularizers and optimizers are not significantly found to be improved.